

Credit Card Fraud Detection with XGBoost

The background of the slide is a dark blue gradient. It features a complex network of glowing blue lines and nodes, resembling a data visualization or a neural network. The lines are thin and connect various points, some of which are highlighted with larger, brighter blue circles. The overall effect is a sense of digital connectivity and data flow.

Samalka Anandagoda

Problem and summary of the Data


Problem:

Predict whether a credit card transaction is fraudulent or not.

Estimate the net savings by identifying fraudulent activity.


Data Overview:

- 10,000 credit card transactions.
- 10 features including transaction_id, amount, transaction_hour, foreign_transaction, location mismatch, etc.
- Target variable: is_fraud = 0 or 1.
- Highly unbalanced dataset.
 - 98.49% non-fraudulent transactions
 - 1.51% fraudulent transactions.




Exploratory Data Analysis (EDA)

Transaction density for every hour



- More fraudulent transactions are taking place during the early hours.
- Non-fraudulent transactions are evenly distributed across the hours.

Categorical Features




- Fraud counts are lower than the non-fraud counts as expected.
- **Merchant category** -> doesn't show a significant variation in fraud activities based on the category.
- **Foreign transaction** -> Fraud counts are similar for both foreign and local transactions. Non-fraud counts follow a different trend
- **Location mismatch** -> Fraud counts are similar for both cases of location mismatch and match. Non-Fraud counts is high for location_mismatch = 0 and low for location_mismatch = 1

Location mismatch and Foreign transactions seem to be good features to classify fraud from non-fraud transactions

Exploratory Data Analysis (EDA)

Feature Density Plots



- The following features appear to be good predictors to classify fraudulent activities from non-fraudulent activities.

Transaction_hour,
device_trust_score,
location_mismatch,
foreign_transaction,
velocity_last_24h

Building model using XGBoost

XGBoost (eXtreme Gradient Boosting):

- Fast and efficient
- Excellent for imbalanced datasets which is the case here
- Built-in regularization to avoid overfitting

Accuracy metric used: **ROC-AUC** (Area under the Receiver Operating Characteristic Curve) which is excellent for testing accuracy for imbalanced datasets.

model Test AUC score: 99.99%

Cross validation score:

obtained a cross-val score of 99.98%

Final Results:

- Model built using XGBoost can classify non-fraudulent from fraudulent transactions with **an accuracy metric (AUC score) of 0.9998 (99.98%)**
- Based on the Feature Importance, location_mismatch, velocity_last_24h, foreign_transaction are effective predictors in classifying.
- Using the ML model, a total **monthly value of 58,638 \$** was saved by identifying fraudulent transactions.

